

AMATH 482 Homework 4

Jessica Shi, 1762333

March 6, 2021

Abstract

In this homework, we use the Linear Discriminant Analysis(LDA) to analyze the MNIST data set and do classifications for data with different digit labels. We also use the Support Vector Machines(SVM) and decision tree classifiers and compare the result of classifying on the same data set.

1 Introduction and Overview

In this homework, I have the MNIST data set with training data, testing data and labels. I would first perform an analysis on the data by doing an SVD analysis and project the data into PCA space. Also, I would plot the related information about the singular values and plot the projection onto 3 selected V-modes. Then, I would implement several classifiers to identify data with different labels. I would first use the Linear Discriminant Analysis to classify between 2 labels. Then, similarly, by changing the old method, I would perform LDA to classify 3 different labels. I would try to find the pair of digits that are most difficult and easiest to separate. Besides, I would use the Support Vector Machines and the decision tree classifiers to do the same classification on the data. And I would compare the performance of each classifier on both training and testing sets.

2 Theoretical Background

The fundamental concept used in this homework is the Principal Component Analysis and also the Linear Discriminant Analysis. And they are both derived method based on the Singular Value Decomposition.

2.1 Principal Component Analysis

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* \quad (1)$$

This is the SVD of a matrix \mathbf{A} . Note that \mathbf{V}^* is the transpose of \mathbf{V} . We would use this to implement the Principal Component Analysis. We introduce the covariance matrix.

$$\mathbf{A} = \frac{1}{\sqrt{n-1}}\mathbf{X} \quad (2)$$

$$\mathbf{C}_\mathbf{X} = \frac{1}{n-1}\mathbf{X}\mathbf{X}^* = \mathbf{A}\mathbf{A}^* \quad (3)$$

Also, we know that

$$\mathbf{C}_\mathbf{X} = \mathbf{A}\mathbf{A}^* = \mathbf{U}\mathbf{\Sigma}^2\mathbf{U}^* \quad (4)$$

So here, the eigenvalues of the covariance matrix are the squares of the scaled singular values. And recall that we used a change of basis to work in the basis of principal components. To do this, I multiply the data by $\mathbf{U}^{-1} = \mathbf{U}^*$:

$$\mathbf{Y} = \mathbf{U}^{-1}\mathbf{X} \quad (5)$$

The covariance \mathbf{Y} is:

$$\mathbf{C}_\mathbf{Y} = \frac{1}{n-1}\mathbf{Y}\mathbf{Y}^T = \frac{1}{n-1}\mathbf{U}^T\mathbf{X}\mathbf{X}^T\mathbf{U} = \mathbf{U}^T\mathbf{A}\mathbf{A}^T\mathbf{U} = \mathbf{U}^T\mathbf{U}\mathbf{\Sigma}^2\mathbf{U}^T\mathbf{U} = \mathbf{\Sigma}^2 \quad (6)$$

In the actual implementation, I would use the builtin function `svd(A)` to get the corresponding \mathbf{U} , $\mathbf{\Sigma}$ and \mathbf{V}^* matrices for \mathbf{A} .

2.2 Linear Discriminant Analysis

The Linear Discriminant Analysis(LDA) helps us to find a proper chosen subspace to project our data sets on in order to have a clear separation between them. The goal is to find a projection that maximized the distance between the inter-class data while minimizing the intra-class data.

Then we define the between-class scatter matrix:

$$\mathbf{S}_B = (\mu_2 - \mu_1)(\mu_2 - \mu_1)^T \quad (7)$$

and also the within-class scatter matrix:

$$\mathbf{S}_w = \sum_{j=1}^2 \sum_x (x - \mu_j)(x - \mu_j)^T \quad (8)$$

Our goal is to find a vector \mathbf{w} such that:

$$\mathbf{w} = \underset{\mathbf{w}}{\operatorname{argmax}} \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_w \mathbf{w}} \quad (9)$$

And it turns out that this vector is the eigenvector corresponding to the largest eigenvalue of the generalized eigenvalue problem:

$$\mathbf{S}_B \mathbf{w} = \lambda \mathbf{S}_w \mathbf{w} \quad (10)$$

So, we could actually use MATLAB to solve for this vector and find proper projection.

3 Algorithm Implementation and Development

According to the specification, I will divide the implementation section into 2 parts.

3.1 SVD Analysis of Data

First, I load the data into MATLAB using the given `mnist_parse` function to get both the training data and the test data, and their corresponding labels. Then I reshape each image into a column vector and perform the PCA on the training data. I use the `svd` function to get the \mathbf{U} , $\mathbf{\Sigma}$ and \mathbf{V} matrix. The matrix \mathbf{U} represents the modes themselves. The matrix $\mathbf{\Sigma}$ gives us the singular value to calculate the energy in each mode and to determine how many modes are necessary for image reconstruction. And the columns in \mathbf{V} represents projections onto each mode. Lastly, I plot the projection onto three selected \mathbf{V} -modes and color the data by their digit label.

3.2 Classifier Implementation

After the basic analysis of the data, I start to implement the classifiers.

First, I implement the linear classifier for 2 classes using LDA. I use the code from lecture 19 and write a function `trainer`. And it works properly for classifying between 2 different labels. Here, I take the label 4 and label 5 as examples. I also write the function `getData` to extract the data with certain labels. After the training process, I plot the histogram for each label and draw a vertical line for the decision-making threshold. From the `trainer` function, I get the proper \mathbf{U} matrix so that I could use it for the projection of test data. Then I apply this matrix to test data and find the final accuracy for the testing process.

Then, I start to implement the linear classifier for 3 classes. Basically, the idea is similar to the 2-class classifier. So, I revise the function `trainer` to get a proper classifier for this. I write a new function

`trainer3` to classify among 3 different labels and I also have the function `getThreshold` to get the 2 different thresholds between these 3 groups. Also, I use the function `getData3` to get the data for the 3 selected labels.

Then I use the builtin decision tree classifier and the SVM to separate the data according to the labels as well. I first get the models using the training data and then use the model to get the predicted labels for test data. Next, I calculate the accuracy for both the training and testing process. The detail for the implementation is in the Appendix B.

4 Computational Results

4.1 SVD Analysis Result

By doing the SVD analysis of the data, I get the following plot for the singular value and the corresponding energies.




Figure 1: Singular Value and Energy

I then calculate the rank for the digit space. I set the threshold for energy to be 90% and I get a rank of 87. So we could say that 87 modes are needed for the image reconstruction for this data set.

Also, I plot the projection onto 3 selected V-modes, which are column 2, 3 and 5. And I use different colors to represent the data with different labels. Below is the plot.




Figure 2: Projection on V-modes

4.2 Classification Result

First, I choose the digit 4 and 5 to do one-time LDA on them. I plot the histogram of the classification result on the train data. And I also calculate the training accuracy, which is 98.37% for label 4 and 98.23% for label 5. We could also see from the Figure 3 that the accuracy should be high.

Then, by applying the \mathbf{U} matrix to the test data, I also do the classification on test data and get an accuracy of 47.6%. And I loop over all the digits to try different combinations and I found that the highest accuracy appears in the pair of label 0 and label 1, which is 73.00%. And the lowest accuracy appears in the pair of label 2 and label 3, which is 35.31%. This might mean that pair (0,1) would be the easiest to separate while pair (2,3) would be the most difficult to separate.

For the 3 classes classifier, I choose the digit 4, 5 and 9. I plot the classification result histogram, which is shown in Figure 4. While we get a relatively good result for digit 5 and 9, we get undesired result for digit 4 and 9. The accuracy is relatively low. The accuracy on test data is 34.06%. This is reasonable since it is hard to use a linear classifier for 3 classes.

Next, I use the builtin function of decision tree classifier and use the function `predict` to get the prediction of labels for test data and check the accuracy. The highest accuracy appears in the pair of label 1 and label 2, which is 15.60%. And the lowest accuracy appears in the pair of label 0 and label 9, which is 82.40%. This means that pair (1,2) would be the most difficult to separate while pair (0,9) would be the easiest to separate.

Lastly, I use the Support Vector Machine(SVM). By similar approach, first I train on the training set and then use the resulting model for the test data. By looping over all the combinations of digits, I get the accuracy table. According to the table, the highest accuracy appears in the pair of label 1 and label 9, which

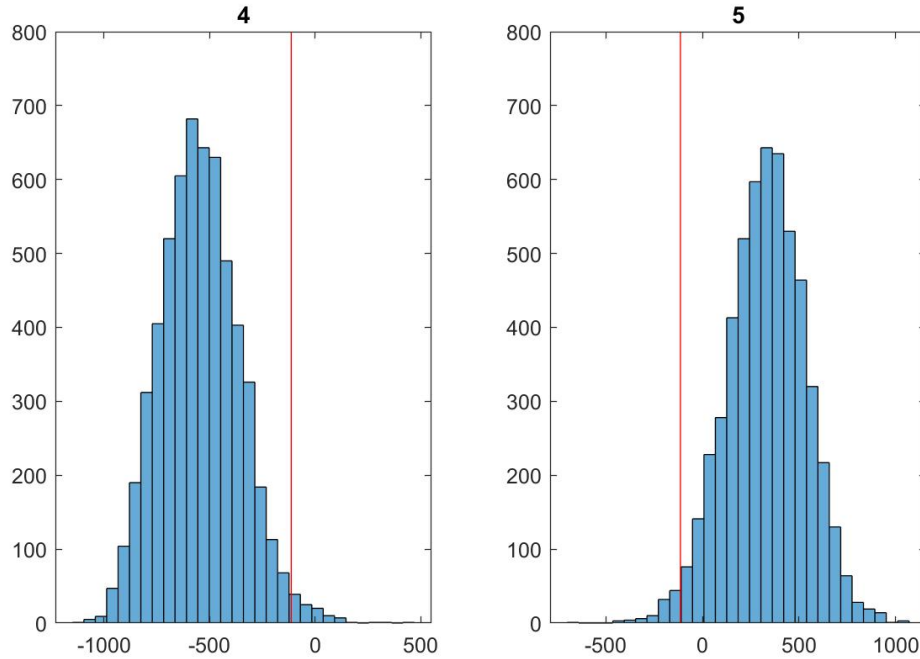


Figure 3: Classification of Train Data between label 4 and 5

is 78.64%. And the lowest accuracy appears in the pair of label 1 and label 2, which is 20.54%. This might mean that pair (1,9) would be the easiest to separate while pair (1,2) would be the most difficult to separate for this classifier.

5 Summary and Conclusions

To conclude, in order to classify images with different labels, we could use Linear Discriminant Analysis as well as other models such as decision tree and Support Vector Machine. The linear classifier is useful for classifying between 2 groups of data, but might be less useful when we have 3 classes. Also, different classifiers or different models might have various performance on the same data set. From the computational result we know that the pair of digits that are easiest and most difficult to separate are not same for LDA, decision tree and SVM. They all perform well on the training accuracy but get an accuracy of about 40-60% for the test data, which are not as good as the training set. And for different pairs of digits, they have diverse performances.

Appendix A MATLAB Functions

- `x = diag(A)` returns a column vector of the main diagonal elements of `A`.
- `plot(X,Y)` creates a 2-D line plot of the data in `Y` versus the corresponding values in `X`.
- `plot3(X,Y,Z)` plots coordinates in 3-D space.
- `[row,col] = find(...)` returns the row and column subscripts of each nonzero element in array `X` using any of the input arguments in previous syntaxes.
- `[U,S,V] = svd(A)` performs a singular value decomposition of matrix `A`, such that $A = U \cdot S \cdot V'$.




Figure 4: Classification of Train Data between label 4, 5, 9

- `tree = fitctree(Tbl,Y)` returns a fitted binary classification decision tree based on the input variables contained in the table `Tbl` and output in vector `Y`.
- `Mdl = fitcsvm(Tbl,Y)` returns an SVM classifier trained using the predictor variables in the table `Tbl` and the class labels in vector `Y`.
- `label = predict(SVMModel,X)` returns a vector of predicted class labels for the predictor data in the table or matrix `X`, based on the trained support vector machine (SVM) classification model `SVMModel`. The trained SVM model can either be full or compact.
- `loss = kfoldLoss(CVMdl)` returns the classification loss obtained by the cross-validated, binary kernel model (ClassificationPartitionedKernel) `CVMdl`. For every fold, `kfoldLoss` computes the classification loss for validation-fold observations using a model trained on training-fold observations. By default, `kfoldLoss` returns the classification error.

Appendix B MATLAB Code

```

1 clear all; close all; clc
2
3 %% load the training data
4 % [images, labels] = mnist_parse('train-images-idx3-ubyte', 'train-labels-idx1-ubyte');
5 [train_data, train_label] = mnist_parse('train-images-idx3-ubyte', 'train-labels-idx1-ubyte');
6 [test_data, test_label] = mnist_parse('t10k-images-idx3-ubyte', 't10k-labels-idx1-ubyte');
```

```

7
8 % SVD analysis
9 % reshape the train data: each column is an image
10 train_data = double(reshape(train_data, size(train_data,1)*size(train_data,2), []));
11 train_label = double(train_label);
12 test_data = double(reshape(test_data, size(test_data,1)*size(test_data,2), []));
13 test_label = double(test_label);
14
15 % use PCA
16 [M, N] = size(train_data);
17 train_data = train_data - repmat(mean(train_data, 2), 1, N);
18 [U, S, V] = svd(train_data/sqrt(N-1), 'econ');
19 % [U, S, V] = svd(train_data, 'econ');
20 sig = diag(S);
21 proj_data = U' * train_data;
22
23 % calculate the rank of the digit space
24 energy = 0;
25 total = sum(diag(S).^2);
26 thre = 0.9;
27 r = 0;
28 while energy < thre
29     r = r+1;
30     energy = energy + (S(r,r).^2)/total;
31 end
32
33 %% plot the singular value and energy
34 figure(1)
35 subplot(2, 1, 1)
36 plot(sig, 'bo', 'Linewidth', 0.5)
37 ylabel('\sigma')
38
39 subplot(2, 1, 2)
40 plot(sig.^2 / sum(sig.^2), 'bo', 'Linewidth', 0.5)
41 ylabel('Energy')
42
43 %% Projection onto 3 V-modes
44 % 2, 3, 5 for now
45 for label=0:9
46     label_indices = find(train_label == label);
47     plot3(V(label_indices, 2), V(label_indices, 3), V(label_indices, 5), ...
48         'o', 'DisplayName', sprintf('%i',label), 'Linewidth', 2)
49     hold on
50 end
51 xlabel('2nd V-Mode'), ylabel('3rd V-Mode'), zlabel('5th V-Mode')
52 title('Projection onto V-modes 2, 3, 5')
53 legend
54 set(gca, 'FontSize', 10)
55
56 % -----
57 %% Implement LDA for 2 digits -- 2 selected digits
58 % get the data with label 4 and 5:
59 l1 = 4;
60 l2 = 5;
61 proj_data = U' * train_data;
62 [data1, data2, n1, n2, label] = getData(l1, l2, train_label, proj_data);
63
64 feature = 87;
65
66 [U,S,V,threshold,w,sort1,sort2] = trainer(data1, data2, feature);
67
68 % calculate the accuracy for train data
69 result1 = sort1 > threshold;
70 result2 = sort2 < threshold;
71 err1 = size(find(result1==1));
72 err2 = size(find(result2==1));
73 acc1 = 1 - err1/size(sort1);
74 acc2 = 1 - err2/size(sort2);

```

```

75
76 % plot the result for train data:
77 figure(2)
78 subplot(1,2,1)
79 histogram(sort1,30); hold on, plot([threshold threshold], [0 800], 'r')
80 %set(gca,'Xlim',[-3 4], 'Ylim',[0 10], 'FontSize',14)
81 title(l1)
82 subplot(1,2,2)
83 histogram(sort2,30); hold on, plot([threshold threshold], [0 800], 'r')
84 %set(gca,'Xlim',[-3 4], 'Ylim',[0 10], 'FontSize',14)
85 title(l2)
86
87 % check the accuracy for test data-----
88 % get the test data with label 4 and 5:
89 [data1, data2, n1, n2, label] = getData(l1, l2, test_label, test_data);
90 test = [data1 data2];
91
92 % change first label to 0, second to 1 to do comparison
93 label = zeros(n1+n2,1);
94 label(1:n1) = 0;
95 label(n1+1:n1+n2) = 1;
96 TestNum = size(test, 2);
97 TestMat = U' * test;
98 pval = w' * TestMat;
99 ResVec = (pval > threshold);
100 % 0's are correct, 1s are incorrect
101 err = abs(ResVec - label');
102 errNum = sum(err);
103 accuracy = 1 - errNum/TestNum;
104 % -----
105 %% Implement LDA for 2 digits: try different combinations:
106 % get the data with label 4 and 5:
107 accu_LDA = zeros(10, 10);
108 for i = 0:9
109     for j = i+1:9
110         l1 = i;
111         l2 = j;
112
113         [data1, data2, n1, n2, label] = getData(l1, l2, train_label, proj_data);
114
115         feature = 87;
116
117         [U,S,V,threshold,w,sort1,sort2] = trainer(data1, data2, feature);
118
119         % calculate the accuracy for train data
120         result1 = sort1 > threshold;
121         result2 = sort2 < threshold;
122         err1 = size(find(result1==1));
123         err2 = size(find(result2==1));
124         acc1 = 1 - err1/size(sort1);
125         acc2 = 1 - err2/size(sort2);
126
127         % check the accuracy for test data-----
128         % get the test data
129         [data1, data2, n1, n2, label] = getData(l1, l2, test_label, test_data);
130         test = [data1 data2];
131
132         % change first label to 0, second to 1 to do comparison
133         label = zeros(n1+n2,1);
134         label(1:n1) = 0;
135         label(n1+1:n1+n2) = 1;
136         TestNum = size(test, 2);
137         TestMat = U' * test;
138         pval = w' * TestMat;
139         ResVec = (pval > threshold);
140         % 0's are correct, 1s are incorrect
141         err = abs(ResVec - label');
142         errNum = sum(err);

```



```

143         accuracy = 1 - errNum/TestNum;
144         accuLDA(i+1,j+1) = accuracy;
145     end
146 end
147
148 %% Implement LDA for 3 digits: (select 4,5,9 as example)
149 l1 = 4;
150 l2 = 5;
151 l3 = 9;
152 [data1, data2, data3, n1, n2, n3, label] = getData3(l1, l2, l3, train_label, proj_data);
153 feature = 87;
154
155 [U, S, V, w, sort1, sort2, sort3] = trainer3(data1, data2, data3, feature);
156
157 threshold1 = getThreshold(sort2, sort3);
158 threshold2 = getThreshold(sort3, sort1);
159
160
161 % plot the result for train data:
162 figure(2)
163 subplot(2,2,1)
164 histogram(sort1,30); hold on, plot([threshold2 threshold2], [0 800], 'r')
165 title(l1)
166 subplot(2,2,2)
167 histogram(sort3,30); hold on, plot([threshold2 threshold2], [0 800], 'r')
168 title(l3)
169 subplot(2,2,3)
170 histogram(sort3,30); hold on, plot([threshold1 threshold1], [0 800], 'r')
171 title(l3)
172 subplot(2,2,4)
173 histogram(sort2,30); hold on, plot([threshold1 threshold1], [0 800], 'r')
174 title(l2)
175
176 % check the accuracy for test data-----
177 % get the test data with label 4 and 5:
178 [data1, data2, data3, n1, n2, n3, label] = getData3(l1, l2, l3, test_label, test_data);
179 test = [data1 data2 data3];
180
181 % change first label to 0, second to 1 to do comparison
182 TestNum = size(test, 2);
183 TestMat = U' * test;
184 pval = w' * TestMat;
185 ResVec = zeros(1, TestNum);
186 for i = 1:TestNum
187     if pval(i) > threshold2
188         ResVec(i) = l1; % label1
189     elseif pval(i) < threshold1
190         ResVec(i) = l2; % label2
191     else
192         ResVec(i) = l3;
193     end
194 end
195 err_num = 0;
196 for i = 1:TestNum
197     if (ResVec(i) != label(i))
198         err_num = err_num + 1;
199     end
200 end
201 sucRate = 1 - err_num/TestNum;
202
203 %% decision tree: different combination of digits
204 % choose data with labels 4 & 5:
205 accuDecTree = zeros(10,10);
206 for i = 0:9
207     for j = i+1:9
208         l1 = i;
209         l2 = j;
210         [data1, data2, n1, n2, label] = getData(l1, l2, train_label, proj_data);

```

```

211     data = [data1 data2];
212     tree = fitctree(data', label, 'CrossVal', 'On');
213     % view(tree.Trained{1}, 'Mode', 'graph');
214     classError = kfoldLoss(tree);
215
216     [data1, data2, n1, n2, label] = getData(l1, l2, test_label, test_data);
217     test = [data1 data2];
218     test_labels = predict(tree.Trained{1}, test');
219     accu = sum(test_labels == label)/length(label);
220     accu_DecTree(i+1, j+1) = accu;
221 end
222 end
223
224 %% SVM classifier: different combination of digits
225 % multiplying SV by the inverse of the largest/first singular value.
226 % choose data with labels 4 & 5:
227 accu_SVM = zeros(10,10);
228 for i = 0:9
229     for j = i+1:9
230         l1 = i;
231         l2 = j;
232         [data1, data2, n1, n2, label] = getData(l1, l2, train_label, proj_data);
233         data = [data1 data2];
234         data = data ./ max(data(:));
235         Mdl = fitcsvm(data', label);
236         [data1, data2, n1, n2, label] = getData(l1, l2, test_label, test_data);
237         test = [data1 data2];
238         test_labels = predict(Mdl, test');
239
240         % 0 false; 1 true
241         check = label==test_labels;
242         err = size(find(check==0), 1);
243         accu = 1 - err/size(test_labels, 1);
244         accu_SVM(i+1,j+1) = accu;
245     end
246 end

```

```

1 function [data1, data2, n1, n2, label] = getData(label1, label2, labels, data)
2     ind1 = find(labels==label1);
3     ind2 = find(labels==label2);
4     n1 = size(ind1, 1);
5     n2 = size(ind2, 1);
6     data1 = zeros(784, n1);
7     data2 = zeros(784, n2);
8     for i = 1:n1
9         ind = ind1(i);
10        data1(:,i) = data(:, ind);
11    end
12
13    for i = 1:n2
14        ind = ind2(i);
15        data2(:,i) = data(:, ind);
16    end
17    label = zeros(n1+n2,1);
18    label(1:n1) = label1;
19    label(n1+1:n1+n2) = label2;
20
21 end

```

```

1 function [data1, data2, data3, n1, n2, n3, label] = getData3(label1, label2, label3, ...
2     labels, data)
3     ind1 = find(labels==label1);
4     ind2 = find(labels==label2);
5     ind3 = find(labels==label3);

```

```

5     n1 = size(ind1, 1);
6     n2 = size(ind2, 1);
7     n3 = size(ind3, 1);
8     data1 = zeros(784, n1);
9     data2 = zeros(784, n2);
10    data3 = zeros(784, n3);
11    for i = 1:n1
12        ind = ind1(i);
13        data1(:,i) = data(:, ind);
14    end
15
16    for i = 1:n2
17        ind = ind2(i);
18        data2(:,i) = data(:, ind);
19    end
20
21    for i = 1:n3
22        ind = ind3(i);
23        data3(:,i) = data(:, ind);
24    end
25    label = zeros(n1+n2+n3,1);
26    label(1:n1) = label1;
27    label(n1+1:n1+n2) = label2;
28    label(n1+n2+1:n1+n2+n3) = label3;
29
30 end

```

```

1 function [U, S, V, threshold, w, sort1, sort2] = trainer(data1, data2, feature)
2
3     n1 = size(data1, 2);
4     n2 = size(data2, 2);
5     [U, S, V] = svd([data1 data2], 'econ');
6     % projection onto principal componenet: X = USV' -> U'X = SV'
7     proj = S*V';
8     U = U(:, 1:feature);
9     d1 = proj(1:feature, 1:n1);
10    d2 = proj(1:feature, n1+1:n1+n2);
11    m1 = mean(d1, 2);
12    m2 = mean(d2, 2);
13
14    Sw = 0;
15
16    for k = 1:n1
17        Sw = Sw + (d1(:,k)-m1)*(d1(:,k) - m1)';
18    end
19    for k = 1:n2
20        Sw = Sw + (d2(:,k)-m2)*(d2(:,k) - m2)';
21    end
22    Sb = (m1 - m2)*(m1 - m2)';
23
24    [V2, D] = eig(Sb, Sw);
25    [lambda, ind] = max(abs(diag(D)));
26    w = V2(:, ind);
27    w = w/norm(w,2);
28    v1 = w' * d1;
29    v2 = w' * d2;
30
31    if mean(v1) > mean(v2)
32        w = -w;
33        v1 = -v1;
34        v2 = -v2;
35    end
36
37    sort1 = sort(v1);
38    sort2 = sort(v2);
39    t1 = length(sort1);

```

```

40     t2 = 1;
41     while sort1(t1) > sort2(t2)
42         t1 = t1-1;
43         t2 = t2+1;
44     end
45
46     threshold = (sort1(t1) + sort2(t2))/2;
47 end

```

```

1 function [U, S, V, w, sort1, sort2, sort3] = trainer3(data1, data2, data3, feature)
2
3     n1 = size(data1, 2);
4     n2 = size(data2, 2);
5     n3 = size(data3, 2);
6     [U, S, V] = svd([data1 data2 data3], 'econ');
7     % projection onto principal componenet: X = USV' -> U'X = SV'
8     proj = S*V';
9     U = U(:, 1:feature);
10    d1 = proj(1:feature, 1:n1);
11    d2 = proj(1:feature, n1+1:n1+n2);
12    d3 = proj(1:feature, n1+n2+1:n1+n2+n3);
13    m1 = mean(d1, 2);
14    m2 = mean(d2, 2);
15    m3 = mean(d3, 2);
16    mu = mean([m1 m2 m3], 2);
17
18    % Calculate within-class matrix
19    Sw = 0;
20    for k = 1:n1
21        Sw = Sw + (d1(:,k)-m1)*(d1(:,k) - m1)';
22    end
23    for k = 1:n2
24        Sw = Sw + (d2(:,k)-m2)*(d2(:,k) - m2)';
25    end
26    for k = 1:n3
27        Sw = Sw + (d3(:,k)-m3)*(d3(:,k) - m3)';
28    end
29
30    % Calculate between-class matrix
31    Sb = (mu - m1)*(mu - m1)' + (mu - m2)*(mu - m2)' + (mu - m3)*(mu - m3)';
32
33    [V2, D] = eig(Sb, Sw);
34    [lambda, ind] = max(abs(diag(D)));
35    w = V2(:, ind);
36    w = w/norm(w,2);
37    v1 = w' * d1;
38    v2 = w' * d2;
39    v3 = w' * d3;
40
41    sort1 = sort(v1);
42    sort2 = sort(v2);
43    sort3 = sort(v3);
44
45 end

```

```

1 function threshold = getThreshold(sort1, sort2)
2     t1 = length(sort1);
3     t2 = 1;
4     while sort1(t1) > sort2(t2)
5         t1 = t1-1;
6         t2 = t2+1;
7     end
8     threshold = (sort1(t1) + sort2(t2))/2;
9 end

```

```

1 function [images, labels] = mnist_parse(path_to_digits, path_to_labels)
2
3 % The function is courtesy of stackoverflow user rayryeng from Sept. 20,
4 % 2016. Link: ...
5     https://stackoverflow.com/questions/39580926/how-do-i-load-in-the-mnist-digits-and-label-data-in-matlab
6 % Open files
7 fid1 = fopen(path_to_digits, 'r');
8
9 % The labels file
10 fid2 = fopen(path_to_labels, 'r');
11
12 % Read in magic numbers for both files
13 A = fread(fid1, 1, 'uint32');
14 magicNumber1 = swapbytes(uint32(A)); % Should be 2051
15 fprintf('Magic Number - Images: %d\n', magicNumber1);
16
17 A = fread(fid2, 1, 'uint32');
18 magicNumber2 = swapbytes(uint32(A)); % Should be 2049
19 fprintf('Magic Number - Labels: %d\n', magicNumber2);
20
21 % Read in total number of images
22 % Ensure that this number matches with the labels file
23 A = fread(fid1, 1, 'uint32');
24 totalImages = swapbytes(uint32(A));
25 A = fread(fid2, 1, 'uint32');
26 if totalImages ≠ swapbytes(uint32(A))
27     error('Total number of images read from images and labels files are not the same');
28 end
29 fprintf('Total number of images: %d\n', totalImages);
30
31 % Read in number of rows
32 A = fread(fid1, 1, 'uint32');
33 numRows = swapbytes(uint32(A));
34
35 % Read in number of columns
36 A = fread(fid1, 1, 'uint32');
37 numCols = swapbytes(uint32(A));
38
39 fprintf('Dimensions of each digit: %d x %d\n', numRows, numCols);
40
41 % For each image, store into an individual slice
42 images = zeros(numRows, numCols, totalImages, 'uint8');
43 for k = 1 : totalImages
44     % Read in numRows*numCols pixels at a time
45     A = fread(fid1, numRows*numCols, 'uint8');
46
47     % Reshape so that it becomes a matrix
48     % We are actually reading this in column major format
49     % so we need to transpose this at the end
50     images(:, :, k) = reshape(uint8(A), numCols, numRows).';
51 end
52
53 % Read in the labels
54 labels = fread(fid2, totalImages, 'uint8');
55
56 % Close the files
57 fclose(fid1);
58 fclose(fid2);
59
60 end

```